

Problem. Five published halting signals stop a reasoning model early. Does any of them generalise across benchmarks and models?

1. Evidence

six signals in twelve settings, one protocol that calibrates, then tests
one rule spans +8.3 to -23.5 points

2. Mechanism

wrong answers repeat more often as the answer space shrinks
stickiness predicts where agreement fails

3. Controller

one shared probe stream; tiers ordered by what they require
worst-case regret 0.031, against 0.100

Guarantee, stress-tested. Holds in 97% of cells within a domain; the worst transferred cell loses 50 points; required in every domain, it certifies only the policy that never halts.

Outcome. Certify halting per domain, or deploy the fusion tier chosen to be safe across domains.