

Reducing the cost of inference in a reasoning language model

```
graph TD; A[Reducing the cost of inference in a reasoning language model] --> B[Smaller model<br/>distillation,<br/>quantisation, pruning<br/>orthogonal:<br/>cheaper tokens]; A --> C[Shorter trace<br/>a fixed budget, or halting<br/>during generation<br/>this dissertation]; A --> D[Cheaper model<br/>route easy queries<br/>to a small model<br/>complementary: joint tier]; A --> E[Faster serving<br/>batching, paged attention,<br/>speculative decoding<br/>orthogonal:<br/>cheaper tokens];
```

Smaller model

distillation,
quantisation, pruning

orthogonal:
cheaper tokens

Shorter trace

a fixed budget, or halting
during generation

this dissertation

Cheaper model

route easy queries
to a small model

complementary: joint tier

Faster serving

batching, paged attention,
speculative decoding

orthogonal:
cheaper tokens